


Train time per epoch (s)



Inference latency (ms/batch)

